

MOHAMMAD MUBASHIRUDDIN

Senior Consultant | Big Data Administrator (Cloudera Admin) | L3 Hadoop

Parbhani, Maharashtra, India · +91 96611 18055 · er.md.mubashir@gmail.com

PROFESSIONAL SUMMARY

Big Data Administrator with 10+ years of hands-on experience implementing and managing Cloudera platforms (CDP, Cloudera Manager, CDH, Hortonworks/Ambaari) across on-premise and AWS-hosted environments. Currently the senior escalation point for four large production CDP clusters at NatWest, covering cluster management, performance tuning, Cloudera Manager upgrades (7.11.3 to 7.13.1) and patching, capacity planning, replication and disaster recovery, security controls (Ranger, Kerberos, Active Directory, KMS) and vendor management. Experienced in source system integration (Sqoop, Flume, Talend-orchestrated pipelines), containerisation with Docker and Kubernetes, CI/CD tooling (Jenkins, Git, Ansible, Terraform) and ITIL-aligned operations including on-call, change management and daily platform reports.

CORE SKILLS

Cloudera & Big Data Platform (L3): CDP Private Cloud, Cloudera Manager, CDH, Hortonworks/Ambaari, Cloudera Support case management, cluster installation and implementation, upgrades, patching, commissioning/decommissioning, performance tuning, HA configuration

Capacity, Backup & DR: Capacity planning and management, Cirata / WANdisco (LiveData replication, cross-environment troubleshooting), cluster-to-cluster data and code migration, Hive backup and restore validation, HA configuration

Security & Compliance: Apache Ranger policies, Kerberos, Active Directory integration, keytab management, KMS key rotation, patching, auditing and compliance controls

Source System Integration: Sqoop (RDBMS to HDFS), Flume, Oozie, Hue workflows, Talend-orchestrated pipeline support

Hadoop Ecosystem: HDFS, YARN, Hive (tuning, query debugging, metastore/warehouse issues), Impala, Hue, Oozie, Tez, Sqoop, Flume, Pig, Spark workloads

Cloud, Containers & CI/CD: AWS (EC2, VPC, subnets, security groups, IAM, CloudFormation, snapshots, S3), Cloudera Manager deployment on AWS, Docker, Kubernetes (pods, services, basic cluster operations), Ansible, Terraform, Jenkins, Git

Monitoring, ITSM & Reporting: Grafana, Ambaari, Resource Manager, Cloudera Manager alerts and events, ITIL incident / change / service request process, on-call and major incident handling, daily platform reports

OS & Domains: Linux (RHEL/CentOS), Windows. Banking & Financial Services, Supply Chain, Healthcare, Telecom

PROFESSIONAL EXPERIENCE

Wipro Limited — Senior Consultant (L3 Hadoop Administrator)

May 2025 – Present

Client: NatWest | Project: DataLake Platform

- Administer four large on-premise CDP Hadoop clusters, owning stability, performance and high availability for a regulated banking data platform.
- Led the Cloudera Manager upgrade from 7.11.3 to 7.13.1 across all four production CDP clusters, covering pre-upgrade checks, change approval, execution and post-upgrade validation.
- Act as L3 escalation point for incidents that clear L1/L2, driving root-cause analysis, permanent fixes and post-incident documentation.
- Manage Cirata/WANdisco replication services end to end, including KMS key rotations, replication rule setup and cross-environment troubleshooting.
- Own Hive service operations: performance tuning, query debugging, metastore and warehouse issue resolution for business-critical workloads.
- Design and maintain Apache Ranger policies for secure access control, auditing and compliance across the Hadoop ecosystem.
- Manage stakeholders and vendors, including Cloudera Support, the Unix team, the Cirata team and other third parties, for upgrades, defect resolution and environment stability.

- Run routine platform administration: monitoring, performance tuning, patching, capacity planning, cluster health checks and incident handling.
- Containerise supporting applications with Docker and deploy basic workloads on Kubernetes, covering image lifecycle, pods and services.

Capgemini — Hadoop Administrator

Nov 2019 – Mar 2025

Client: HP | Project: HP TERA (shared multi-tenant cluster, ~50% workload weighting)

- Administered three Hortonworks clusters in a shared, multi-tenant environment where HP TERA carried roughly half the total cluster load.
- Managed Hive operations: table creation, backups, statistics maintenance and resolution of failed or stuck jobs.
- Provisioned and governed access to databases, HDFS paths and service accounts through Ranger policies.
- Monitored platform health and job execution using Ambari, Resource Manager, Tez View and Grafana; tracked queue utilisation and rebalanced capacity.
- Troubleshoot job failures across Talend-orchestrated pipelines, configuring Talend context files and working alongside the EDL and Talend platform teams.
- Performed edge node and HDFS trash cleanup to protect cluster storage and namenode performance.

Larsen & Toubro Infotech (LTI) — Hadoop Administrator

Dec 2017 – Oct 2019

Client: Johnson & Johnson | Project: Enterprise Data Lake

- Installed, configured and administered enterprise Cloudera (CDH) clusters across Dev, QA and Production, including POC build-outs.
- Secured clusters with Kerberos and Active Directory integration; created and managed keytabs and user onboarding.
- Executed cluster-to-cluster migrations covering both data and code, plus Hive backups and restore validation.
- Deployed Cloudera Manager and the Hadoop cluster on AWS and administered the supporting components: VPC, EC2, subnets, security groups, route tables, IAM roles, CloudFormation and snapshots.
- Configured YARN queues, set job priorities and identified users generating excessive small files to protect cluster performance.
- Re-ran and recovered workflows through Hue, identifying stuck jobs and resolving underlying causes.
- Worked strictly within ITIL process (service requests, incidents, change requests), participated in the on-call rota and produced daily platform reports.
- Configured alerts and events, managed hosts, roles, services and backups, and commissioned/decommissioned nodes as capacity demanded.

Defacto Veritas — Hadoop Administrator

Jun 2015 – Nov 2017

Domain: Telecom | Platform support, from POC to production

- Installed and configured open-source Apache Hadoop and Cloudera Manager clusters from proof of concept through to production support.
- Managed HDFS, services, user permissions and High Availability configuration.
- Installed and configured Hive, Pig, Flume and Sqoop; managed RDBMS-to-HDFS data transfer via Sqoop.
- Monitored logs and provided platform support to multiple consuming teams.

EDUCATION

- B.E. Computer Science & Engineering — P.E.S. College of Engineering, Aurangabad (First Division)
- HSC — Queen's Junior College, Parbhani (First Class, 67.67%)
- SSC — Queen's English School, Parbhani (First Class, 74%)

ADDITIONAL DETAILS

Languages: English, Hindi, Marathi (speak, read, write)

Nationality: Indian · Valid Indian passport (valid until December 2028)

Notice Period: Available on request